


The Four Levels of Agentic Maturity

Move from chatbots and co-pilots to autonomous, collaborative agents with a clear maturity model for capability, governance, data, and value.



LEVEL 0

Fixed rules & repetitive tasks

LEVEL 1

Information retrieval agents

LEVEL 2

Simple orchestration, single domain

LEVEL 3

Complex orchestration, multiple domain

LEVEL 4

Multi-agent orchestration



Fixed Rules and Repetitive Tasks

Chatbots and co-pilots are the early stage of agentic AI adoption. The next step is moving toward true agents that can operate autonomously and collaborate with humans.

What it is

Automation of repetitive tasks using predefined rules, with no reasoning or learning capabilities. At this level, chatbots and co-pilots are primarily focused on retrieving information, such as answering FAQs or handling basic inquiries. They help automate tasks, but do not recommend actions or take further steps. The shift to agentic AI begins when an agent moves beyond simply retrieving information to using reasoning to recommend or take action autonomously, as seen in higher levels of maturity.

Example

Chatbots handling basic service inquiries, like password resets or FAQ responses, following simple decision trees without adapting to new information.

HOW TO ADVANCE FROM LEVEL 0 TO LEVEL 1

Identify use cases where chatbots and co-pilots are limited by rigid decision trees, and explore opportunities for reasoning capabilities to improve outcomes.

Conduct time studies to quantify capacity saved by having agents automate tasks and recommend actions.

Select use cases aligned with the organisation's risk tolerance and establish a mitigation strategy — for example, employee-facing versus customer-facing agents in regulated industries.

Ensure data quality and availability by starting with a harmonized data source and plan for a broader data strategy as maturity increases.

Focus not only on time savings but also on value and personalized customer experiences.

Quantify time savings and demonstrate early efficiency metrics, such as reduced response times and cost savings from automation.



Information Retrieval Agents

Agents that assist humans by retrieving information and recommending actions.

What it is

Agents that assist humans by retrieving information and recommending actions.

Example

An agent pulling data from a knowledge base to suggest the next steps for a customer support ticket, such as recommending relevant troubleshooting articles or escalating the issue.

HOW TO ADVANCE FROM LEVEL 1 TO LEVEL 2

Shift from recommendations to actions, allowing agents to autonomously execute tasks, reducing manual error and speeding up completion.

Harmonize additional data, ensuring it is cleansed, transformed, and pre-processed as new sources are integrated.

Establish a broader governance framework with testing methods, performance metrics, and continuous monitoring.

Consider how user feedback will be gathered and incorporated.

For customer-facing agents, ensure brand consistency is reflected.

Measure ROI based on efficiency and effectiveness metrics like time savings, case deflection rates, and customer satisfaction.



Simple Orchestration, Single Domain

Agents that autonomously orchestrate low-complexity tasks in a siloed data environment.

What it is

Agents that autonomously orchestrate low-complexity tasks in a siloed data environment.

Example

An agent that schedules meetings and automates follow-up emails using data from internal calendar and email systems.

HOW TO ADVANCE FROM LEVEL 2 TO LEVEL 3

Decide whether to build one multi-functional agent or multiple agents working together to achieve common goals. Consider latency impacts when choosing between a mega-agent or specialized agents.

Grant agents access similar to human workers, ensuring separation of duties and applying the principle of least privilege.

Design architecture to support future complexity, incorporating standards and protocols to avoid constant re-architecture.

Use out-of-the-box API connectors to extend agents' capabilities beyond single domains.



Complex Orchestration, Multiple Domain

Agents that autonomously orchestrate multiple workflows with harmonized data across multiple domains.

What it is

Agents that autonomously orchestrate multiple workflows with harmonized data across multiple domains.

Example

An agent handling complex cross-department workflows, such as managing the sales pipeline while pulling data from CRM, customer service tickets, and financial reports to create a holistic customer view.

HOW TO ADVANCE FROM LEVEL 3 TO LEVEL 4

Focus on use cases that require real-time, cross-domain agent collaboration, optimizing complex workflows with dynamic agent teams.

Develop a universal agent communication layer — APIs or an “agent bus” — for seamless collaboration.

Implement dynamic agent discovery, including registration, de-registration, and de-duplication, for ongoing maintenance.

Design a scalable architecture that supports any-to-any agent interactions.

3

Governance, Supervision & Value

As agents operate across domains, controls must mature alongside their autonomy and reach.

Establish fine-grained access control, security policies, and a robust governance framework covering auditability, error handling, and transparency.

Create layered human/AI supervision based on a risk framework.

Develop a sustainable AI agent lifecycle management methodology.

Measure ROI by assessing efficiency, effectiveness, experience, and risk mitigation.

MATURITY LEVEL 4

Multi-Agent Orchestration

Any-to-any-agent operability across disparate stacks with agent supervision.

Example

Multiple agents from different systems collaborating to autonomously process orders, manage inventory, and route customer feedback across various departments in real time.

HOW TO MAXIMIZE VALUE

Evolve security and governance for ecosystem-wide collaboration.

Explore new business models enabled by agent collaboration.

Develop metrics to quantify multi-agent system value.


Create feedback mechanisms to improve agent interactions.

Measure ROI by tracking overall business impact from multi-agent collaboration, including revenue growth, cost savings, and customer retention.

BEYOND THE TECHNOLOGY

Critical Considerations for Agent Implementation

Organisations must look beyond the technology and consider the broader organisational impact — including data readiness, security, and the need for a collaborative human-agent workforce.



AI agents require more than technical implementation. They demand a strong foundation in strategy, data, security, and workforce readiness. Organisations must align AI initiatives with business goals, ensure high-quality and integrated data sources, establish governance frameworks as agent autonomy increases, and invest in upskilling teams to foster seamless human-agent collaboration. Addressing these pillars accelerates adoption and drives sustainable AI maturity.